

2026 Global Agent Competition - Track B Judging

Rubric

(Culture, Commerce, Tourism, and Sports Skills Solution Package)

I. General Rules and Scoring Method

1. Scoring Framework (five dimensions; 100 points total):

- Technical Implementation and Engineering Quality (Skills Package) - 25 points: package completeness, required number of Skills, and quality of individual Skills.
- Application Value and Industry Solution (Pitch) - 25 points: real business value, quality of the industry solution, and a closed-loop business model.
- Expert / Representative Enterprise Knowledge Distillation - 20 points: distilling and packaging industry know-how from public materials.
- Innovation and User Experience - 15 points: instruction quality, usability, and originality.
- Presentation and Demo - 15 points: demonstration video.

2. Scoring Principles (evidence-based):

- Each score must be supported by Skills Package execution logs, test-case results, observed outputs, or knowledge-distillation materials. Scores may not be based on impressions alone.
- This rubric does not set “featured/non-featured” thresholds. More complete and substantive materials and solutions receive higher scores. The benchmarks in this rubric define the full-score range; scores decrease

according to coverage when a benchmark is not met and may be as low as zero.

- Dimensions are scored independently: strength in one dimension does not offset a deficiency in another.
- Use the full scoring range to distinguish quality and avoid score clustering. Comparable issues should receive comparable scores across evaluation batches.

3. Final Score: The final composite score for Track B is the average of all expert judges' scores.

4. Track Positioning and Award-Tier Notes:

- This track is open to schools, companies, and OPC developers (individuals or teams of 1–5 people). Using the prescribed theme of Culture, Commerce, Tourism, and Sports (culture; business and trade; tourism and cultural tourism; sports), entrants must build a Skills Package that addresses a real business need of a representative company in the relevant industry.
- Expert scoring reflects only the quality of the entry. To receive a First Prize or another high-tier award recognized by the Organizing Committee, an entry must meet the applicable score threshold and all High-Tier Award Eligibility Requirements. Fewer than six valid Skills fails the entry gate and is not scored. An entry that passes the gate but does not meet the high-tier requirements for engineering complexity, depth of distillation, or Demo may receive no higher than Second Prize.

II. Required Submission Materials (Final Submission)

Prescribed theme: Culture, Commerce, Tourism, and Sports (culture; business and trade; tourism and cultural tourism; sports). Each team must select at least one sub-area and build and deliver a Skills Package for a real business need of a

representative enterprise, such as a tourism group, mixed-use commercial complex, sports club, or cultural institution.

Each team may submit one Skills Package entry. The final submission must include:

- 1) Skills Package source package (required; pre-submission validation): at least six valid Skills. Each Skill must include a SKILL.md file, scripts or logic, test cases, and resources. The system will automatically verify completeness and publication status. A valid Skill has a defined role/function, can be invoked independently with correct output, contains substantive business logic, and is not merely a templated wrapper.
- 2) Industry solution brief (PDF; required; pitch deck source): target-company profile, real business pain points and needs, mapping between the Skills Package and business touchpoints, an end-to-end business loop, and measurable metrics for cost reduction, efficiency gains, or experience improvements.
- 3) Expert / representative-enterprise knowledge-distillation package (required): compiled public materials, an explanation of how public industry knowledge or enterprise SOPs are packaged into Skills, and a source list for domain knowledge.
- 4) Demonstration environment (required): a workspace or agent link bound to the Skills Package that supports online execution and invocation. If only a ZIP file is submitted, the post-import runtime environment and an authentic screen recording will be used as demonstration evidence.
- Demo materials (recommended): an MP4 screen recording of no more than five minutes as a backup for the live demonstration. A live link that can be invoked in real time is preferred; ZIP submissions must be imported live or accompanied by an authentic runtime recording.
- 6) API documentation (PDF; required): API calls and parameter descriptions for each Skill in the Skills Package.

- 7) Entry cover and summary (required): a cover image and a summary of no more than 3000 characters describing the theme, target company, Skills Package overview, and core value.

III. Compliance and Security Entry Gate (Automatic Disqualification)

The following are mandatory preconditions and are outside the 100-point score. If any condition applies, the entry is deemed ineligible, will not proceed to dimension scoring, and is disqualified from awards. Integrity and security violations—including malicious code, data-security violations, privilege bypass, unlawful content, falsified Skills, authorizations, or test cases, and fabricated distillation materials—cannot be cured through remediation. Non-integrity technical issues, such as the number of valid Skills, file format, broken links, or ZIP compatibility, may be corrected and re-reviewed before the submission deadline as directed by the Organizing Committee.

- Security red line: malicious code such as viruses, cryptomining software, backdoors, or exploits.
- Data security: theft or exfiltration of user data, keys, credentials, or other sensitive information.
- Access-control compliance: bypassing system security mechanisms, unauthorized access, or privilege escalation.
- Content compliance: generating or distributing unlawful content, pornography, violence, or politically sensitive content.
- Behavioral consistency: actual behavior differs from the stated functionality, or hidden functions are present.
- Supply-chain security: introducing third-party dependency versions with known vulnerabilities.
- Falsified authorizations or test cases.

- The Skills Package contains fewer than six valid Skills, lacks a ZIP file or another valid independently importable validation artifact, or contains falsified Skills.
- Fabricated distillation materials: falsified public sources, altered industry data, or invented industry rules or enterprise SOPs.
- The entry does not use any Culture, Commerce, Tourism, or Sports sub-area as its theme, or lacks an identifiable representative enterprise or real business scenario (incomplete entry).
- Reusing knowledge-base materials, imagery, or Q&A content across agents, workspaces, or tasks.

IV. High-Tier Award Eligibility Requirements (Award Thresholds; Not Automatic Disqualification)

For First Prize (and, if established separately, Champion, Runner-up, Third Place, or similar honors), an entry must meet the applicable expert-score threshold and all of the following requirements:

- The Skills Package meets at least four Engineering Depth and Complexity Baseline (Cx) requirements: (1) at least one real tool, API, or MCP call—not text generation alone; (2) a workflow of at least three steps with exception branches; (3) at least one domain-specific rule that matches the industry model and is demonstrably effective; (4) robustness, with clear fallback behavior or errors for at least three types of invalid input; and (5) nontrivial integration that forms a business loop with an agent or workspace.
- Expert / representative-enterprise distillation is complete: the authenticity of the distillation and depth of knowledge packaging meet the standard, with verifiable supporting materials.
- Live Demo requirement: explanation plus live demonstration totals at least 10 out of 12 points.
- The package contains at least six valid Skills.

Interpretation: If most Skills in a package are “minimal atomic Skills” (only one interface, no business orchestration, and direct pass-through of inputs and outputs), the package is deemed not to meet the complexity requirement. Regardless of the expert score, it may receive no higher than Second Prize and may not compete for First Prize or, if established separately, Champion, Runner-up, Third Place, or similar honors. This clause does not affect the entry’s original expert score; it limits only the award tier.

Dimension 1: Technical Implementation and Engineering Quality (Skills Package) (25 points)

ID	Scoring Item	Points	What Judges Observe	Full-Score Benchmark (Rubric Threshold)	If Not Met / Scoring Scale
1.1	Package completeness and required quantity	10	Does the Skills Package contain at least six valid Skills? Does it cover at least three business functions (for example, consulting, booking, guiding, marketing, or data)? Are all Skills published to the Skills marketplace or independently verifiable as a package? A valid Skill includes SKILL.md, scripts/logic, and test cases; can be invoked independently with correct	At least six valid Skills; at least three functional categories; all runnable (10 points).	Six or more Skills but limited functional breadth or inconsistent runtime quality (6–8 points). Fewer than six valid Skills, inability to validate independently, or falsified Skills: 0 for 1.1, triggers the entry gate, and no further scoring.

			output; and contains substantive business logic rather than a template wrapper.		
1.2	Functional validity of individual Skills	8	Randomly test at least three Skills: test cases cover normal, boundary, and error conditions; outputs are correct; formats are consistent; and fallback behavior is graceful.	All sampled Skills meet the reliability baseline (8 points); approximately 2.7 points per item.	Minus 0.5 for each missing test category; minus 1 per factual error; minus 1.5 for a silent failure; minus 1 for unstable formatting.
1.3	Engineering standards and integration	7	Complete SKILL.md, script, and resource directories; effective workspace/agent integration; and performance and resource efficiency.	Complete directory structure (3); working integration (2); acceptable performance (2).	Minus 0.5 per missing directory item; 0 if integration cannot work. A template-only solution or one that only changes an API endpoint is capped at 3 points.

Dimension 2: Application Value and Industry Solution (Pitch)

(25 points)

ID	Scoring Item	Points	What Judges Observe	Full-Score Benchmark (Rubric Threshold)	If Not Met / Scoring Scale
2.1	Quality of the industry application	12	Is the solution brief clear about the target-company profile,	All five elements are present, with measurable metrics (12	Minus 2 for no company profile; minus 3 for fabricated pain

	solution		real pain points/needs, mapping of the Skills Package to business touchpoints, end-to-end business loop, and measurable metrics (cost reduction, efficiency, and experience)? It serves as the pitch source.	points).	points; minus 3 for no business loop; minus 2 for no measurable metrics.
2.2	Scenario value and problem solving	8	Do the Skills address real, high-frequency business needs and strongly match Culture, Commerce, Tourism, or Sports scenarios? Both function description and triggering scenario are required.	Strong value (8 points); addresses real, high-frequency needs.	Minus 3 for a weak scenario; 0 if unrelated to the scenario. A general information query with no business-process improvement is capped at 4 points.
2.3	Business loop and adoption value	5	Practical deployability, reusability, industry fit, and a closed-loop business model.	Complete (5 points).	Minus 2 if only partly achieved; 0 if not practical.

Dimension 3: Expert / Representative Enterprise Knowledge Distillation (20 points)

ID	Scoring Item	Points	What Judges Observe	Full-Score Benchmark (Rubric	If Not Met / Scoring Scale
----	--------------	--------	---------------------	------------------------------	----------------------------

				Threshold)	
3.1	Authenticity of distillation	8	Provide compiled public materials, a knowledge-distillation explanation, and a source list. Materials must be verifiable and traceable.	Detailed public materials, clear sources, and correspondence with Skill logic (8 points).	Thin materials or incomplete source traceability (4 points). Falsified or altered sources: 0 and triggers the entry gate. No materials: 0.
3.2	Depth of knowledge packaging	7	Does expert experience or enterprise SOPs genuinely enter Skill logic through a knowledge base, decision branches, or business rules, rather than being cited by name only? Sampled Skills should demonstrate industry know-how.	Deep packaging (7 points).	Surface-level citation (3 points); no packaging (0 points).
3.3	Domain-model fit	5	Does the entry match the relevant industry model (cultural tourism, culture, commerce, or sports) and meet sector-specific dimensions such as cultural authenticity, local accuracy, visitor experience, trade compliance, or sports rules?	Matches the model and meets the standard (5 points).	Partial match (3 points); no match (0 points).

Dimension 4: Innovation and User Experience (15 points)

ID	Scoring Item	Points	What Judges Observe	Full-Score Benchmark (Rubric Threshold)	If Not Met / Scoring Scale
4.1	Instruction quality and innovative design	5	Clarity and executability of instructions, error handling, and boundary coverage.	Clear, imperative, example-driven; includes prechecks and exception branches (5 points).	Minus 1 per ambiguous statement; minus 1.5 for no prechecks; minus 1.5 for no exception branch.
4.2	User experience and usability	5	Output-format consistency, clarity of SKILL.md guidance, and usability for developers and users.	Stable formatting, clear guidance, and usable out of the box (5 points).	Requires users to guess how to use it (3 points); disorganized and not independently usable (0 points).
4.3	Technical or pattern novelty	5	Novel capability, reusability, and industry-model fit. Changing only an endpoint or rewriting prompts does not constitute meaningful differentiation.	Clear differentiator (5 points): introduces a capability or industry adaptation not covered by the template and makes it work in practice.	General template only (2 points); copied template without adaptation (0 points).

Dimension 5: Presentation and Demo (15 points)

ID	Scoring Item	Points	What Judges Observe	Full-Score Benchmark (Rubric Threshold)	If Not Met / Scoring Scale
5.1	Industry solution presentation	7	Are the industry-solution logic, value proposition,	Complete and clear (7 points).	Minus 3 for partial explanation; 0 if

			and business loop explained clearly?		key elements are missing.
5.2	Demo demonstration	5	Does the demo invoke the core capabilities of the Skills Package, including multi-Skill integration, with verifiable outputs? Use of AI tools to build the demonstration live is encouraged.	Complete and verifiable (5 points).	Minus 2 for a partial demo; 0 if a key capability fails.
5.3	Communication	3	Presentation logic.	Clear and accurate (3 points).	Minus 2 if unclear.

V. Score Bands and Use of Results

Award mapping rules (expert scoring represents entry quality; final awards in the national review also depend on composite ranking and compliance with the High-Tier Award Eligibility Requirements):

- If Champion, Runner-up, and Third Place are established separately: top composite rankings, final expert score of at least 85, and satisfaction of all High-Tier Award Eligibility Requirements.
- First Prize: final expert score of at least 80 and satisfaction of all High-Tier Award Eligibility Requirements. An entry that fails any high-tier eligibility requirement may not receive First Prize.
- Second Prize: final expert score of at least 70. An entry that passes the entry gate but does not meet the high-tier requirements for engineering complexity, depth of distillation, or Live Demo may receive no higher than Second Prize.
- Third Prize and Special Awards: final expert score generally at least 60 and satisfaction of all entry-gate requirements. A Special Award also requires

verifiable evidence in the applicable area of technology, commercial implementation, or public engagement.

- Below 60: Ineligible; may not receive any competition award.

Note: Score bands reflect only expert evaluation of entry quality. Final rankings for the domestic Track B competition are based on the final composite score. Award eligibility is also subject to the entry gate and High-Tier Award Eligibility Requirements in this rubric.

Rating	Range (Expert Score)	Description
Excellent	90 or above	Near-perfect, with only minor defects; suitable as a model for demonstration and wider adoption.
Good	80–89	Strong overall quality, with a small number of areas for improvement.
Fair	70–79	Meets the basics but has notable shortcomings that need refinement.
Needs Improvement	60–69	Core functions are usable, but multiple standards are not met and substantial iteration is needed.
Ineligible	Below 60	Serious deficiencies or an entry-gate violation; disqualified from awards.