

From a request to organized work

Skill function description | SP-D Organizational Collaboration Office | 27 September 2026

A workspace that makes collaboration inspectable

Atlas Office brings programmable staff profiles, task routing, skill execution, review gates and generated business files into one local workspace. A colleague has a manager, assigned skills, availability, capacity, follow links and role guidance. The reviewer can inspect who performed a task, what it produced and what still needs attention.

The entry combines a general company workspace with a specialized Chengdu tourism demonstration. Its 69 registered skills include nine independently callable tourism stages. Local execution creates real draft documents and structured records; the people, supplier offers and tourism task evidence are synthetic.

At a glance

69 registered skills	9 tourism stages	4 interfaces
Office tasks, analysis, governance and optional connected actions	A composable inquiry-to-outcome simulation	Browser workspace, HTTP, CLI and MCP

Submission scenario: Product Launch Preparation

The supplied SP-D challenge specifies Product Launch Preparation. The narrated recording follows one request through an agent-led launch simulation for a Chengdu food and tea group-tour product. The recorded run completed 20 skill executions across 16 distinct skills and 7 fictional colleagues, producing 10 real files. It used 24 model requests and 26 tool calls; reported usage totaled 800,795 tokens across requests. Actual wall time was 303.8 seconds.

The agent compares pilot and full-launch options, checks team capacity, records decisions, schedules dependencies, maps responsibility, prepares a campaign and launch checklist, tracks tasks, assesses risk, calculates budget variance and a finance forecast, creates the executive pack, and saves and retrieves reusable knowledge. The pack includes editable Word, PowerPoint and Excel files plus a PDF report.

Who it serves

A sales or operations lead preparing a group-travel product, proposal or internal launch decision; a specialist creating schedules and comparisons; and a manager reviewing blockers, ownership and release conditions. Public information about Sichuan Shanghai Airlines Holiday International Travel Agency Co., Ltd. provides a representative travel-business context, without implying adoption or affiliation.

Project repository

<https://github.com/energelpen/1010-Chengdu-Hackathon>

Sources: README.md; resources/company-general.json; resources/tourism-company.json | Challenge Topics.pdf, physical page 6; source-register.md

Nine skills, one tourism lifecycle

Independently callable stages

Stage / skill	Function and observable result
01 inquiry-intake	Parse or validate group size, dates, duration and CNY budget; preserve supplied versus extracted fields; return missing-input issues.
02 flight-search	Filter fixture offers by route, dates, seats and price validity; compare group totals and preserve unconfirmed-fare caveats.
03 tour-search	Rank activities by stated interests and check capacity and accessibility; explain rejected candidates.
04 option-comparison	Construct complete quotation alternatives with itemized cost, margin and budget checks; rank feasible options.
05 hierarchy-router	Select a lead and duty owners using skills, availability and capacity; expose capability gaps and reporting routes.
06 workload-splitter	Create owned tasks with dependencies, preparation dates and approval gates for supplier, safety, finance and management.
07 schedule-builder	Place selected activities into dated slots with arrival/departure and timing constraints; surface unscheduled items.
08 execution-controller	Accept a completion event only with the assigned owner, evidence, dependencies and gate conditions; record blockers and escalations.
09 outcome-reporter	Assemble management status, evidence and a customer email draft; hold release when required checks remain incomplete.

A reproducible synthetic input

RFQ: 30 travelers for a 3 day Chengdu food, tea and heritage tour from 2026-10-25, budget CNY 200000.

Decision rules carry the business logic

Group capacity and the customer budget constrain the quotation. Missing supplier, safety or finance evidence blocks release. The demonstration model adds an executive task at CNY 100,000. These thresholds and approvals are authored demo rules, not the reference enterprise's SOP. Flight and tour options are estimates; execution records do not constitute reservations.

Sources: scripts/tourism_core.py; skills/*/SKILL.md; shared/message-schema.json | resources/domain-rules.md; resources/tourism-price-model.json

Prepare a launch with visible outputs

A concrete launch preparation task

The recording prepares a fictional Chengdu food and tea group-tour launch. A CNY 50,000 pilot is compared with a CNY 120,000 full launch. The agent supplies clearly labeled synthetic assumptions and delegates the actual scoring and proposal generation to proposal-package. Its recommendation and detailed inputs are retained in the execution evidence.

Composable tools for the preparation team

Skill	Business use / output
proposal-package	Compare explicitly supplied cost, impact and risk; generate an editable PowerPoint proposal and Word decision brief.
project-plan	Turn structured tasks and dependencies into a dated project plan; reject invalid dependency input.
responsibility-matrix	Validate colleagues and produce a workbook identifying responsible, accountable, consulted and informed people.
workload-rebalance	Allocate work to qualified staff within capacity; expose assignments that need manager attention.
scenario-compare	Score options with named criteria, weights and visible component scores; reject missing measurements.
management-handoff	Summarize progress, owners and blockers in a Word brief and unsent email draft.
task-tracker / decision-log	Persist task status and decision rationale for subsequent inspection.
quotation-package / finance-forecast	Create a draft commercial quotation and price schedule, or a scenario-based monthly forecast.

Interaction and assignment

The Assistant routes recognized requests to a qualified colleague and records progress. The Skills library offers structured input forms and examples for deliberate invocation. Activity & approvals exposes runs and pending actions; Files & knowledge retains generated artifacts. The Org chart distinguishes reporting, current collaboration and follows.

The recording uses the configured OpenAI model to orchestrate the entire local workflow from one request. Offline general-company chat provides directory/help responses and does not execute free-form requests. Staff profiles supply role context, assigned skills and capacity; one orchestrator performs the delegated tool calls. They are not separate authenticated employee accounts or independent model sessions.

Current integration boundary

One user brief starts a real configured OpenAI model. The model chooses the skills, publishes a dependency-aware plan, assigns qualified colleagues, supplies the inputs and executes local work through the shared runtime. The recording contains no manual skill submissions. All business facts and launch outcomes are explicitly simulated; no external delivery, spending or real product release occurs.

The live dashboard persists the model-authored plan, assignees, prerequisite states, each actual tool call, run IDs, artifact counts, elapsed time and provider-reported token usage. Step completion is based on runtime status. Pending approvals and unfinished plans are never counted as completed execution.

Sources: scripts/intent_router.py; scripts/assistant_service.py; scripts/skill_runtime.py | skills/proposal-package/skill.json; skills/project-plan/skill.json; tests/test_office_skills.py

Architecture, review and traceability

Execution path

Layer	Responsibility
Workspace	Assistant, colleague profiles, organization view, skill forms, run history and files.
Routing and orchestration	Interpret a supported request, choose a qualified person, prepare inputs, invoke the registered handler and expose progress.
Shared skill runtime	Validate input schemas; create run records; apply idempotency and approval states; register artifacts and audit events.
Domain handlers	Tourism decision rules, office analysis and report generation, or a configured connector adapter.
Persistence	Local SQLite stores for runs, conversations and records; an FTS5 case index; generated files in the workspace data directory.

Review is part of the workflow

- The tourism workspace defaults to preparing a proposal for review. Return, revise, resubmit and escalate actions preserve a management history.
- Changing a selected option or task owner invalidates affected approvals and evidence. Task completion cannot bypass proposal approval.
- General skills with `remote_write` or `reviewed_write` effects stop at `awaiting_approval`. Approval executes the prepared action; rejection cancels it.
- SMTP delivery and connected services require configuration and explicit review rules. A generated email draft remains unsent in this submission.

Inspectable results

A shared-runtime result contains the requested `skill_id`, `person_id`, `status`, `inputs`, `output` and `timestamps`. Tourism stage data adds `request_id`, `status`, `issues` and `trace` references. A completed runtime handler may still contain a blocked domain result; callers must inspect both levels.

A known trace metadata defect affects the inner skill label of five tourism handlers. The requested route and outer `skill_id` identify the invoked stage correctly. The API reference discloses this limitation; test success is not presented as proof that every metadata field is correct.

Operating boundary

This is a single-user local prototype bound to 127.0.0.1. Production multi-user use needs identity, authorization, protected audit access and a suitable HTTPS server. Optional supplier, Google, Telegram, SMTP and model services require their own configuration; no live supplier reservation adapter is claimed.

Sources: `scripts/skill_runtime.py`; `scripts/governance.py`; `scripts/conversation_store.py` | `scripts/search_store.py`; `scripts/mcp_bridge.py`; `app.py`

Evidence and reproducibility

Checks rerun for this submission on 27 September 2026

Evidence	Observed result	What it establishes
Tourism behavior suite	180 / 180 passed	20 decision cases per tourism stage; normal, boundary and error conditions.
Automated suite	51 / 51 passed	Integration, assistant, governance and office-skill behaviors checked by the existing suite.
Fixed routing corpus	43 / 43 action prompts; 5 / 5 non-action prompts	Correct first-pass routing on this predefined corpus; not arbitrary-language success.
Manifest validation	69 manifests and examples	Each registered skill has an input schema, an embedded schema-valid example, SKILL.md and a CLI runner.
Runtime API example	CNY 1,500; 15% variance	budget-variance executed in isolated local state with budget 10,000 and actual 11,500.

Reproduce the core checks

```
# In chengdu-tourism-office with dependencies installed
python scripts/tourism_office.py test
python -m unittest discover -s tests -p "test_*.py" -q
python tests/trigger_eval.py

# A single independently callable tourism skill
python scripts/tourism_office.py run inquiry-intake '@examples/rfq.json'

# Start the local workspace with fresh state
python app.py --data-dir ./submission-demo-state
```

Interpretation of the evidence

Tests were rerun locally using the repository’s existing checks. Their captured logs are preserved with the submission sources. External-service operations are mocked or require separate credentials; no real mail, booking or payment was performed to create these materials. The 180 behavior cases are reported separately from the 51 automated tests.

The recorded run completed 20 skill executions across 16 distinct skills and 7 fictional colleagues, producing 10 real files. It used 24 model requests and 26 tool calls; reported usage totaled 800,795 tokens across requests. Actual wall time was 303.8 seconds. The successful run establishes this one-request scenario. It does not establish arbitrary-language reliability, enterprise adoption or measured productivity savings.

What a production pilot should establish

An authorized operator should validate the company’s actual SOPs, replace demo rules and supplier fixtures, and measure preparation time, review rework, task ownership, evidence completeness and exception handling. Proposed pilot measures appear in the Enterprise challenge fit statement; they are targets rather than achieved outcomes.

Sources: submission/source/verification/*.txt; tests/test_skills.py; tests/trigger_eval.py | tests/test_integration.py; tests/test_governance.py; tests/test_office_skills.py